

Practical Work 3 - Regularised Regression Methods

Adib Habbou - Alae Khidour

07-11-2022

Data Analysis

Data Import

```
diabetes_data <- read.table(file = "diabetes.txt", header = TRUE)
```

Data Conversion

```
YBin <- as.numeric(diabetes_data$Y > median(diabetes_data$Y))  
diabetes_data <- diabetes_data[, -11]  
diabetes_data <- cbind(diabetes_data, YBin)
```

```
head(diabetes_data)
```


	AGE	SEX	BMI	BP	S1	S2	S3	S4	S5	S6	YBin
1	59	2	32.1	101	157	93.2	38	4	4.8598	87	1
2	48	1	21.6	87	183	103.2	70	3	3.8918	69	0
3	72	2	30.5	93	156	93.6	41	4	4.6728	85	1
4	24	1	25.3	84	198	131.4	40	5	4.8903	89	1
5	50	1	23.0	101	192	125.4	52	4	4.2905	80	0
6	23	1	22.6	89	139	64.8	61	2	4.1897	68	0

```
tail(diabetes_data)
```

	AGE	SEX	BMI	BP	S1	S2	S3	S4	S5	S6	YBin
437	33	1	19.5	80.00	171	85.4	75	2.00	3.9703	80	0
438	60	2	28.2	112.00	185	113.8	42	4.00	4.9836	93	1
439	47	2	24.9	75.00	225	166.0	42	5.00	4.4427	102	0
440	60	2	24.9	99.67	162	106.6	43	3.77	4.1271	95	0
441	36	1	30.0	95.00	201	125.2	42	4.79	5.1299	85	1
442	36	1	19.6	71.00	250	133.2	97	3.00	4.5951	92	0

Data Visualization

```
pairs(diabetes_data, pch = 22,
      bg = c("tomato1", "royalblue3")[unclass(factor(diabetes_data[, "YBin"])]))
```




Red squares correspond to observations where $YBin$ is equal to 0 which means Y is lower than the median. Blue squares correspond to observations where $YBin$ is equal to 1 which means Y is greater than the median. From the above plot, we observe the following:

- Colinearity between $S1$ and $S2$
- Above a certain value for BMI and BP we only find blue squares

Study of correlation

```
corr <- cor(diabetes_data)
corrplot(corr, method = "circle")
```



The correlation plot provides us with a lot of information such as:

- S1 and S2 are highly positively correlated;
- S3 and S4 are highly negatively correlated.

We need to keep in mind the correlation between our variables.

The potential colinearity between variables can have an impact on the Standard Error.

More than that, it means that the co-variable signifacitivity test is useless.

Logistic Regression

Data Partitionning

```
sample <- sample(c(TRUE, FALSE), nrow(diabetes_data), replace = TRUE, prob = c(0.8, 0.2))
train_data <- diabetes_data[sample, ]
test_data <- diabetes_data[!sample, ]
```

Application of logistic regression

```
reg_log = glm(formula = YBin ~ ., family = binomial, data = train_data)
summary(reg_log)
```

Call:

```
glm(formula = YBin ~ ., family = binomial, data = train_data)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.4429	-0.7059	-0.1318	0.7664	2.2188

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-7.9900967	4.7479201	-1.683	0.092402 .
AGE	0.0056811	0.0117240	0.485	0.627980
SEX	-1.1953437	0.3356120	-3.562	0.000368 ***
BMI	0.1408921	0.0396128	3.557	0.000375 ***
BP	0.0284984	0.0119963	2.376	0.017521 *
S1	0.0112509	0.0474511	0.237	0.812576
S2	-0.0235857	0.0471034	-0.501	0.616567
S3	-0.0755260	0.0578409	-1.306	0.191636
S4	-0.0023359	0.3633716	-0.006	0.994871
S5	1.6177768	1.1781280	1.373	0.169698
S6	-0.0009758	0.0153589	-0.064	0.949341

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 482.42 on 347 degrees of freedom
Residual deviance: 322.75 on 337 degrees of freedom
AIC: 344.75

Number of Fisher Scoring iterations: 5

Interpretation of the results

- For the Deviance Residuals we observe that they are close to be centered on 0 and roughly symmetrical.
- We can make assumptions on the significativity of the co-variables by looking at their p-values. For instance *BMI* and *S5* are highly significant co-variables for our model meanwhile *AGE* and *S4* are less significant.
- The dispersion parameter in our case is equal to 1, but we can adjust it if we want too. Since we are not estimating the variance from the data instead we are just deriving it from the mean, it is possible that the variance is underestimated.
- The Akaike Information Criterion (*AIC*) will help us to compare between different models.
- The number of Fisher Scoring iterations tells us how quickly the function converges to the maximum likelihood estimated for the coefficients.

Study of the coefficients

```
reg_log$coefficients
```

(Intercept)	AGE	SEX	BMI	BP
-7.9900966794	0.0056811233	-1.1953437486	0.1408921394	0.0284983749
S1	S2	S3	S4	S5
0.0112508931	-0.0235857284	-0.0755260460	-0.0023359257	1.6177767505
S6				
-0.0009758189				

We should also not only look at the estimated value of the coefficient to determine co-variable significativity because even a very low estimated coefficient can become bigger at the end depending on the co-variable unit and magnitude.

-> The most significant co-variable is *S5*, the less significant co-variable is *AGE*.

But we need to keep in mind that it does not mean that *AGE* is not significant in reality, it is just the least significant in our model according to the computed p-values for our data set.

Predictions type response

```
predict_response <- predict.glm(reg_log, newdata = test_data, type = "response")
predict_response
```

10	12	14	17	18	19	59
0.94368200	0.14816924	0.58901218	0.81300489	0.71755897	0.57865236	0.06534977
60	62	63	65	71	74	76
0.28440924	0.77701566	0.02490906	0.25034394	0.05220693	0.20698489	0.30798449
80	81	82	83	85	97	103
0.11507647	0.58331661	0.23717402	0.03178885	0.06432964	0.89087369	0.34062404
106	109	111	118	124	140	145
0.23206270	0.94596593	0.02101791	0.96525571	0.39774408	0.97925217	0.64308084
150	156	165	167	169	186	194
0.81430630	0.65341847	0.32159309	0.07297506	0.88682831	0.86205940	0.40022993
199	200	201	220	223	225	230
0.50315181	0.85068189	0.10423184	0.53809688	0.51916762	0.15259770	0.08645594
233	238	240	241	242	243	247
0.49627704	0.07009210	0.67049569	0.91694482	0.16042649	0.12016256	0.31884745
248	257	259	260	266	277	278
0.09207751	0.97620816	0.23393675	0.47849027	0.26373068	0.69239687	0.21830626
279	282	287	304	305	308	309
0.28168542	0.16517691	0.04517958	0.94883083	0.20601715	0.43720950	0.15827339
310	318	323	325	331	334	339
0.21463602	0.64275366	0.99658186	0.90834148	0.61381483	0.76013877	0.38356020
340	342	346	356	365	369	370
0.48121388	0.66473506	0.24060972	0.26157232	0.50543566	0.86211128	0.79478631
382	383	386	389	392	403	406
0.11453807	0.99239198	0.35763011	0.62953089	0.06862223	0.52549878	0.98883153
410	411	415	417	419	420	422
0.57710899	0.65776826	0.57545891	0.87120482	0.08253818	0.13413883	0.77392503
423	427	428				
0.63298005	0.50722764	0.27024723				

Using the type *response* we obtain values between 0 and 1 which correspond to the probability of the variable *YBin* being equal to 1 computed from the area under the link function.

Predictions type link

```
predict_link <- predict.glm(reg_log, newdata = test_data, type = "link")
predict_link
```

10	12	14	17	18	19
2.81877508	-1.74903275	0.35988330	1.46965466	0.93238534	0.31724363
59	60	62	63	65	71
-2.66041849	-0.92269424	1.24835897	-3.66729928	-1.09677876	-2.89892103
74	76	80	81	82	83
-1.34319646	-0.80955896	-2.03990435	0.33640343	-1.16823579	-3.41633474
85	97	103	106	109	111
-2.67724281	2.09969660	-0.66051450	-1.19670050	2.86259169	-3.84113830
118	124	140	145	150	156
3.32437769	-0.41487369	3.85434746	0.58876113	1.47823802	0.63409967
165	167	169	186	194	199
-0.74646024	-2.54186278	2.05874535	1.83250102	-0.40450715	0.01260742
200	201	220	223	225	230
1.73995928	-2.15106398	0.15268347	0.07670809	-1.71437054	-2.35769674
233	238	240	241	242	243
-0.01489213	-2.58527543	0.71042786	2.40154214	-1.65505822	-1.99089168
247	248	257	259	260	266
-0.75907351	-2.28852832	3.71433304	-1.18621397	-0.08609207	-1.02666753
277	278	279	282	287	304
0.81134887	-1.27556408	-0.93611669	-1.62020276	-3.05087805	2.92009328
305	308	309	310	318	323
-1.34910247	-0.25249494	-1.67113142	-1.29720361	0.58733598	5.67523573
325	331	334	339	340	342
2.29355045	0.46337635	1.15344048	-0.47446407	-0.07517987	0.68446747
346	356	365	369	370	382
-1.14933966	-1.03781235	0.02174349	1.83293738	1.35402147	-2.04520216
383	386	389	392	403	406
4.87091485	-0.58566506	0.53020486	-2.60804849	0.10208367	4.48342907
410	411	415	417	419	420
0.31091666	0.65336458	0.30415904	1.91165372	-2.40835001	-1.86484926
422	423	427	428		
1.23060832	0.54502269	0.02891255	-0.99336862		

Using the type *link* we obtain the values of the link function.

Odd-Ratios

```
exp(coef(reg_log))
```

(Intercept)	AGE	SEX	BMI	BP	S1
0.0003388013	1.0056972915	0.3025999180	1.1513004614	1.0289083388	1.0113144224
S2	S3	S4	S5	S6	
0.9766902409	0.9272555790	0.9976668005	5.0418685153	0.9990246570	

From the odd-ratios obtained for each co-variable we can evaluate the influence of the co-variable on the target knowing that:

- When the Odd-Ratio is lower than 1 it means that the co-variable had a negative influence on the target, for instance *AGE*, *SEX*, *S1*, *S3*, *S4* and *S6*.
- When the Odd-Ratio is greater than 1 it means that the co-variable had a positive influence on the target, for instance *BMI*, *BP* and *S2*.

The limits of this approach is that if we change the binary labels we had (0 and 1 for our case) we will obtain different values for the estimated coefficients.

Performance

MAP

Using the Maximum A Posteriori criteria we can make predictions for our binary variable *YBin*:

```
prediction <- as.numeric(predict.glm(reg_log, diabetes_data, type = "response") > 0.5)
```

```
table(prediction)
```

prediction	
0	1
226	216

Knowing that our target variable have the following values:

```
table(diabetes_data$YBin)
```

0		1	
221	221		

By comparing the both tables we can tell that our predictions are quite good, since our model has nearly the same count for 0 and 1 than the target variable in our data set.

Confusion Matrix

```
confusion_matrix <- table(diabetes_data$YBin, prediction)
confusion_matrix
```

```
      prediction
      0      1
0 171    50
1   55   166
```

Here we just computed the confusion matrix for our model. This matrix has 4 values which corresponds respectively to the number of True Negative, False Negative, False Positive and True Positive.

- True Negative is the specificity which is the ability to predict $Y\hat{Bin} = 0$ for $YBin = 0$
- True Positive is the sensitivity which is ability to predict $Y\hat{Bin} = 1$ for $YBin = 1$

Accuracy

```
accuracy <- (confusion_matrix[1,1] + confusion_matrix[2,2]) / nrow(diabetes_data)
accuracy
```

```
[1] 0.7624434
```

The accuracy of our model correspond to it's ability to predict the right value (0 or 1) for all the observations of our data set. In our case it's equal to 76 %.

Global Error

```
global_error <- (confusion_matrix[1,2] + confusion_matrix[2,1]) / nrow(diabetes_data)
global_error
```

```
[1] 0.2375566
```

The global error of our model correspond to it's inability to predict the right value (0 or 1) for all the observations of our data set. In our case it's equal to 24 %.

Recall

```
recall <- confusion_matrix[2,2] / (confusion_matrix[1,2] + confusion_matrix[2,2])
recall
```

```
[1] 0.7685185
```

The recall of our model correspond to the correctly predicted positive rate. In our case it's equal to 75 %.

Precision

```
precision <- confusion_matrix[2,2] / (confusion_matrix[2,1] + confusion_matrix[2,2])
precision
```

```
[1] 0.7511312
```

The precision correspond to the rate of correct positive predictions. In our case it's equal to 78 %.

F1-Score

```
f1_score <- (2 * precision * recall) / (precision + recall)
f1_score
```

```
[1] 0.7597254
```

The F1-score correspond to the ability to predict positive individuals well. In our case it's equal to 76 %.

The F_β -score uses a more general formula where β is chosen such that the recall is considered β times as important as the precision:

$$F_\beta = \frac{(1 + \beta^2) * precision * recall}{(\beta^2 * precision) + recall}$$

False Positive Rate

```
confusion_matrix[2,1] / nrow(diabetes_data)
```

```
[1] 0.1244344
```

The false positive is equal to 11 %.

False Negative Rate

```
confusion_matrix[1,2] / nrow(diabetes_data)
```

```
[1] 0.1131222
```

The false negative is equal to 13 %.

K-Fold

The Cross-Validation is a technique which simply reserves a part of the training data and uses it to test the model while the remaining non-reserved data is used to train the model.

The principle behind K-Fold cross validation is that we start by dividing our data set into K equal parts. Then we will train our model on the K-1 first parts and use the last part to test the model. Then we will use another combination of parts to train and test our model until we computed all the possible combinations. In the end, every part of the data set is used for testing and we can then have an idea of the performance of our model on new data.

This technique is used to avoid overfitting and to know the performance of our model on new data.

Here we are coding a function that takes the number of folds and returns a vector of the performance computed at each iteration. The model used here is logistic regression and we are computing the performance by making predictions and computing the confusion matrix

```
kfold_all <- function(k) # k is here the number of folds
{
  # create a vector of length number of folds
  performance <- vector(length = k)

  # create a sequence from 1 to k
  folds <- cut(seq(1,nrow(diabetes_data)), breaks = k, labels = FALSE)

  # perform k fold cross validation
  for(i in 1:k)
  {
    # split data by fold
    index <- which(folds == i, arr.ind = TRUE)
    test_data <- diabetes_data[index,]
    train_data <- diabetes_data[-index,]

    # train the logistic regression on the train data set
    reg_log <- glm(YBin ~ ., family = binomial, data = train_data)

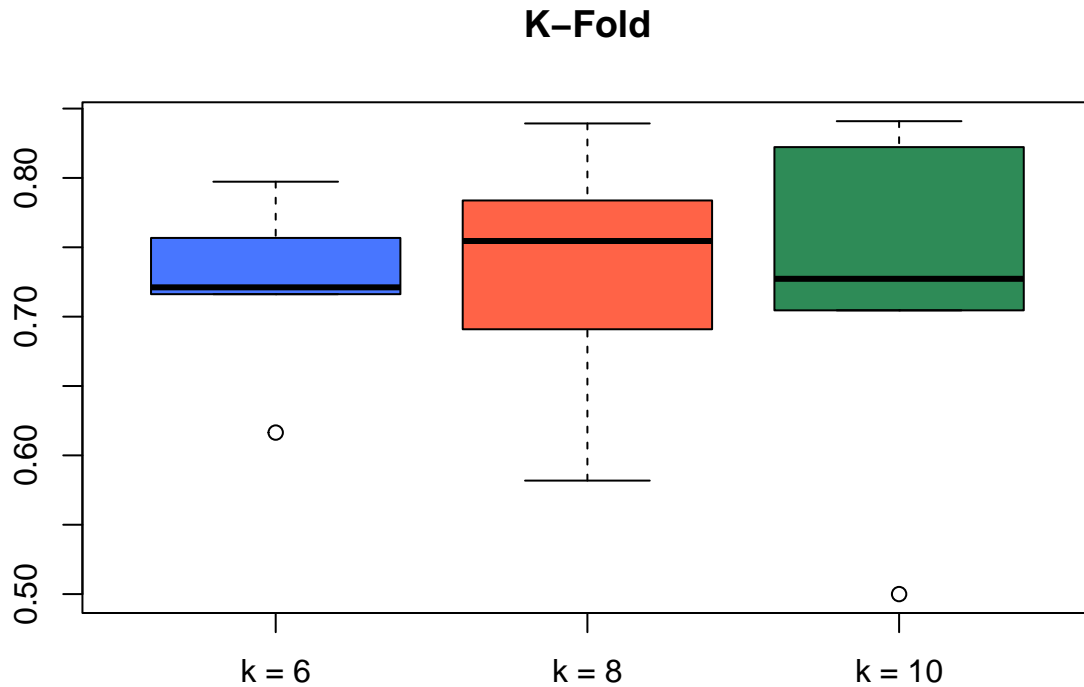
    # make predictions
    prediction <- predict(reg_log, test_data)

    # compute confusion matrix
    confusion_matrix <- table(as.numeric(prediction > 0.5), diabetes_data[index,]$YBin)

    # compute the performance
    performance[i] <- (confusion_matrix[1,1] + confusion_matrix[2,2]) / nrow(test_data)
  }

  # returning the vector of performances
  return (performance)
}
```

```
boxplot(kfold_all(6), kfold_all(8), kfold_all(10),
        main = "K-Fold", names = c("k = 6", "k = 8", "k = 10"),
        col = c("royalblue1", "tomato1", "seagreen"))
```



In our case we did 3 different K-Folds for K equal to 6, 8 and 10. By looking at the boxplots we can see that the median of our accuracy is around 75 % with maximum and minimum values that can go from 60 % to 80 % for the 10-fold.

Variable Selection

Statistical Approach

We start by setting up a model with all the variables and another with only an intercept:

```
reg_all <- glm(YBin ~ ., data = diabetes_data, family = binomial)
reg_none <- glm(YBin ~ 1, data = diabetes_data, family = binomial)
```

Forward Logistic Regression

```
reg_forward <- step(reg_none, list(upper = reg_all), direction = 'forward')
```

Start: AIC=614.74

YBin ~ 1

	Df	Deviance	AIC
+ S5	1	500.61	504.61
+ BMI	1	507.06	511.06
+ BP	1	549.95	553.95
+ S4	1	552.60	556.60
+ S3	1	555.06	559.06
+ S6	1	573.71	577.71
+ S1	1	601.14	605.14
+ AGE	1	601.63	605.63
+ S2	1	604.41	608.41
<none>		612.74	614.74
+ SEX	1	612.73	616.73

Step: AIC=504.61

YBin ~ S5

	Df	Deviance	AIC
+ BMI	1	457.80	463.80
+ BP	1	480.20	486.20
+ S3	1	485.31	491.31
+ S1	1	494.40	500.40
+ SEX	1	497.09	503.09
+ S6	1	497.33	503.33
+ S4	1	497.75	503.75
<none>		500.61	504.61
+ S2	1	499.98	505.98
+ AGE	1	500.18	506.18

Step: AIC=463.8

YBin ~ S5 + BMI

	Df	Deviance	AIC
+ BP	1	448.52	456.52
+ S1	1	449.40	457.40
+ S3	1	450.90	458.90
+ SEX	1	454.15	462.15
+ S2	1	454.55	462.55
<none>		457.80	463.80
+ S4	1	457.56	465.56
+ S6	1	457.56	465.56
+ AGE	1	457.78	465.78

Step: AIC=456.52

YBin ~ S5 + BMI + BP

	Df	Deviance	AIC
--	----	----------	-----

```

+ S1      1    438.95 448.95
+ S3      1    440.93 450.93
+ SEX     1    441.96 451.96
+ S2      1    444.62 454.62
<none>    448.52 456.52
+ AGE     1    448.10 458.10
+ S4      1    448.24 458.24
+ S6      1    448.48 458.48

```

Step: AIC=448.95

YBin ~ S5 + BMI + BP + S1

	Df	Deviance	AIC
+ SEX	1	431.06	443.06
+ S2	1	434.89	446.89
+ S3	1	435.61	447.61
+ S4	1	436.32	448.32
<none>		438.95	448.95
+ AGE	1	438.91	450.91
+ S6	1	438.94	450.94

Step: AIC=443.06

YBin ~ S5 + BMI + BP + S1 + SEX

	Df	Deviance	AIC
+ S2	1	419.13	433.13
+ S3	1	420.41	434.41
+ S4	1	422.48	436.48
<none>		431.06	443.06
+ S6	1	430.85	444.85
+ AGE	1	431.02	445.02

Step: AIC=433.13

YBin ~ S5 + BMI + BP + S1 + SEX + S2

	Df	Deviance	AIC
<none>		419.13	433.13
+ AGE	1	419.00	435.00
+ S4	1	419.12	435.12
+ S3	1	419.12	435.12
+ S6	1	419.13	435.13

We can observe that the forward logistic regression give us a model which contains the following variables: *S5*, *BMI*, *BP*, *S1*, *SEX* and *S5* with an AIC equal to 433.

Backward Logistic Regression

```
reg_back <- step(reg_all, direction = 'backward')
```

Start: AIC=440.97

YBin ~ AGE + SEX + BMI + BP + S1 + S2 + S3 + S4 + S5 + S6

	Df	Deviance	AIC
- S6	1	418.97	438.97
- S3	1	418.97	438.97
- S4	1	418.98	438.98
- AGE	1	419.11	439.11
- S2	1	420.19	440.19
- S1	1	420.79	440.79
<none>		418.97	440.97
- S5	1	427.77	447.77
- BP	1	433.51	453.51
- SEX	1	434.72	454.72
- BMI	1	437.80	457.80

Step: AIC=438.97

YBin ~ AGE + SEX + BMI + BP + S1 + S2 + S3 + S4 + S5

	Df	Deviance	AIC
- S3	1	418.97	436.97
- S4	1	418.98	436.98
- AGE	1	419.12	437.12
- S2	1	420.20	438.20
- S1	1	420.80	438.80
<none>		418.97	438.97
- S5	1	427.88	445.88
- BP	1	434.04	452.04
- SEX	1	434.76	452.76
- BMI	1	438.19	456.19

Step: AIC=436.97

YBin ~ AGE + SEX + BMI + BP + S1 + S2 + S4 + S5

	Df	Deviance	AIC
- S4	1	419.00	435.00
- AGE	1	419.12	435.12
<none>		418.97	436.97
- S2	1	422.28	438.28
- S1	1	427.40	443.40
- BP	1	434.05	450.05
- SEX	1	434.79	450.79
- BMI	1	438.29	454.29
- S5	1	441.41	457.41

Step: AIC=435

YBin ~ AGE + SEX + BMI + BP + S1 + S2 + S5

	Df	Deviance	AIC
--	----	----------	-----

```

- AGE    1    419.13 433.13
<none>      419.00 435.00
- S2     1    431.02 445.02
- BP     1    434.05 448.05
- SEX    1    434.83 448.83
- BMI    1    438.30 452.30
- S1     1    438.91 452.91
- S5     1    480.74 494.74

```

Step: AIC=433.13

YBin ~ SEX + BMI + BP + S1 + S2 + S5

	Df	Deviance	AIC
<none>		419.13	433.13
- S2	1	431.06	443.06
- SEX	1	434.89	446.89
- BP	1	435.39	447.39
- BMI	1	438.47	450.47
- S1	1	438.92	450.92
- S5	1	480.94	492.94

We can observe that the backward logistic regression give us a model which contains the following variables: *S5*, *BMI*, *BP*, *S1*, *SEX* and *S5* with an AIC equal to 433.

Stepwise Logistic Regression

```
reg_both <- step(reg_all, direction = 'both')
```

Start: AIC=440.97

YBin ~ AGE + SEX + BMI + BP + S1 + S2 + S3 + S4 + S5 + S6

	Df	Deviance	AIC
- S6	1	418.97	438.97
- S3	1	418.97	438.97
- S4	1	418.98	438.98
- AGE	1	419.11	439.11
- S2	1	420.19	440.19
- S1	1	420.79	440.79
<none>		418.97	440.97
- S5	1	427.77	447.77
- BP	1	433.51	453.51
- SEX	1	434.72	454.72
- BMI	1	437.80	457.80

Step: AIC=438.97

YBin ~ AGE + SEX + BMI + BP + S1 + S2 + S3 + S4 + S5

	Df	Deviance	AIC
- S3	1	418.97	436.97
- S4	1	418.98	436.98
- AGE	1	419.12	437.12
- S2	1	420.20	438.20
- S1	1	420.80	438.80
<none>		418.97	438.97
+ S6	1	418.97	440.97
- S5	1	427.88	445.88
- BP	1	434.04	452.04
- SEX	1	434.76	452.76
- BMI	1	438.19	456.19

Step: AIC=436.97

YBin ~ AGE + SEX + BMI + BP + S1 + S2 + S4 + S5

	Df	Deviance	AIC
- S4	1	419.00	435.00
- AGE	1	419.12	435.12
<none>		418.97	436.97
- S2	1	422.28	438.28
+ S3	1	418.97	438.97
+ S6	1	418.97	438.97
- S1	1	427.40	443.40
- BP	1	434.05	450.05
- SEX	1	434.79	450.79
- BMI	1	438.29	454.29
- S5	1	441.41	457.41

Step: AIC=435

YBin ~ AGE + SEX + BMI + BP + S1 + S2 + S5

	Df	Deviance	AIC
- AGE	1	419.13	433.13
<none>		419.00	435.00
+ S4	1	418.97	436.97
+ S3	1	418.98	436.98
+ S6	1	419.00	437.00
- S2	1	431.02	445.02
- BP	1	434.05	448.05
- SEX	1	434.83	448.83
- BMI	1	438.30	452.30
- S1	1	438.91	452.91
- S5	1	480.74	494.74

Step: AIC=433.13

YBin ~ SEX + BMI + BP + S1 + S2 + S5

	Df	Deviance	AIC
<none>		419.13	433.13
+ AGE	1	419.00	435.00
+ S4	1	419.12	435.12
+ S3	1	419.12	435.12
+ S6	1	419.13	435.13
- S2	1	431.06	443.06
- SEX	1	434.89	446.89
- BP	1	435.39	447.39
- BMI	1	438.47	450.47
- S1	1	438.92	450.92
- S5	1	480.94	492.94

We can observe that the stepwise logistic regression give us a model which contains the following variables: *S5*, *BMI*, *BP*, *S1*, *SEX* and *S5* with an AIC equal to 433.

Final model

Each time, we obtain the same value for the *AIC* criteria which equals 433. We could choose another criteria for our model selection, such as *BIC* or C_p . It could influence our model because the formula we will minimize changes. This formula will depend differently on the number of co-variables p . So the choice of criteria will always depend on the business constraint behind. More than that, we can customize our criteria as we want by choosing the value of λ in the function *step*.

Our final model will be the following:

```
diabetes_reg <- lm(formula(reg_both), data = diabetes_data)
summary(diabetes_reg)
```

Call:

```
lm(formula = formula(reg_both), data = diabetes_data)
```

Residuals:

Min	1Q	Median	3Q	Max
-1.02548	-0.29301	-0.00166	0.31694	0.94958

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-1.962754	0.188818	-10.395	< 2e-16 ***
SEX	-0.169907	0.042440	-4.003	7.34e-05 ***
BMI	0.024674	0.005261	4.690	3.66e-06 ***
BP	0.006835	0.001605	4.257	2.54e-05 ***
S1	-0.007748	0.001642	-4.719	3.20e-06 ***
S2	0.006534	0.001709	3.823	0.000151 ***
S5	0.457993	0.054360	8.425	5.31e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.4021 on 435 degrees of freedom

Multiple R-squared: 0.3634, Adjusted R-squared: 0.3546

F-statistic: 41.38 on 6 and 435 DF, p-value: < 2.2e-16

From the summary of our regression we can tell that:

- The residuals are quite symmetrically distributed around their median;
- The intercept is equal to -1.962754 , we also notice the influence of each co-variable on Y ;
- The standard error and the t-value are provided to show how the p-values were calculated;
- ALL the p-values are very low which means that all co-variables are significant;
- The R^2 tells us that the p co-variables can explain 36% of the variation in the target variable Y_{Bin} ;
- The first degree of freedom corresponds to $p - 1$ with $p = 7$ the number of variables of the model;
- The second degree of freedom corresponds to $n - p$ with $n = 442$ the number of data points.


Penalized Methods

```
sample <- sample(c(TRUE, FALSE), nrow(diabetes_data), replace = TRUE, prob = c(0.8, 0.2))
train_data <- diabetes_data[sample, ]
test_data <- diabetes_data[!sample, ]
```

Ridge Regression

Regularization Path

```
ridge <- glmnet(x = train_data[,-11], y = train_data$YBin, alpha = 0, family = "binomial")
plot(ridge, xvar = "lambda", label = TRUE, lwd = 2)
```




In the plot above, we can see the evolution of the coefficient values depending on λ . Knowing that we are doing a ridge regression, we can tell that they will all converge to 0.

However, ridge regression does not perform feature selection and will retain all available features in the final model. Therefore, a ridge model is good if we suppose that there is a need to retain all features in our model yet reduce the noise that less influential variables may create.

If greater interpretation is necessary and many of the features are redundant or irrelevant then a lasso or elastic net penalty may be preferable.

Cross-Validation

```
cv_ridge <- cv.glmnet(as.matrix(train_data[,-11]), train_data$YBin,  
                      family = "binomial", alpha = 0, type.measure = "mse")  
plot(cv_ridge)
```



In the above plot, we visualize the cross validation curve:

- λ_{min} is the value which minimizes the Mean Squared Error in Cross-Validation;
- λ_{1se} is the value which is the largest λ value within 1 standard error;
- The intervals estimate variance of the loss metric (red points) using Cross-Validation;
- The vertical lines show the locations of λ_{min} and λ_{1se} ;
- The numbers across the top are the number of non-zero coefficients.

We observe a slight improvement in the Mean Squared Error as our penalty $\log(\lambda)$ gets larger, suggesting that a regular OLS model likely overfits the training data. But as we constrain it further by continuing to increase the penalty of our Mean Squared Error starts to increase.

Lambda Min

```
lambda_min <- cv_ridge$lambda.min
ridge_min <- glmnet(x = train_data[,-11], y = train_data$YBin,
                   alpha = 0, family = "binomial", lambda = lambda_min)
ridge_min$beta
```

```
10 x 1 sparse Matrix of class "dgCMatrix"
      s0
AGE  0.011780226
SEX -1.034321232
BMI  0.144095420
BP   0.033272898
S1  -0.003418393
S2  -0.004901545
S3  -0.031189107
S4   0.096812586
S5   1.299302188
S6   0.007230844
```

The model obtained using λ_{min} gives us the model with lowest Mean Squared Error, which can seem to be a good choice but in fact it can implies overfitting.

```
prediction_min <- as.numeric(
  predict(ridge_min, as.matrix(diabetes_data[,-11]), type = "response") > 0.5)
```

Using the Maximum A Posteriori criteria we can make predictions for our binary variable $YBin$:

```
table(prediction_min)
```

```
prediction_min
 0    1
218 224
```

By comparing with the table of the target variable we can tell that our predictions are quite consistent, since our model has nearly the same count for 0 and 1 than the target variable in our data set.

```
confusion_matrix <- table(diabetes_data$YBin, prediction_min)
confusion_matrix
```

```
      prediction_min
      0    1
0 165  56
1  53 168
```

```
accuracy <- (confusion_matrix[1,1] + confusion_matrix[2,2]) / nrow(diabetes_data)
accuracy
```

```
[1] 0.7533937
```

The accuracy is equal to 76.2%.

Lambda 1se

```
lambda_1se <- cv_ridge$lambda.1se
ridge_1se <- glmnet(x = train_data[,-11], y = train_data$YBin,
                    alpha = 0, family = "binomial", lambda = lambda_1se)
ridge_1se$beta
```

```
10 x 1 sparse Matrix of class "dgCMatrix"
      s0
AGE  0.0068028289
SEX -0.3708382040
BMI  0.0807923072
BP   0.0194120409
S1  -0.0002557427
S2  -0.0015573200
S3  -0.0160877961
S4   0.1067847839
S5   0.6495087342
S6   0.0098241418
```

The model obtained using λ_{1se} gives us a simpler model which can avoid overfitting.

```
prediction_1se <- as.numeric(
  predict(ridge_1se, as.matrix(diabetes_data[,-11]), type = "response") > 0.5)
```

Using the Maximum A Posteriori criteria we can make predictions for our binary variable *YBin*:

```
table(prediction_1se)
```

```
prediction_1se
  0    1
220 222
```

By comparing with the table of the target variable we can tell that our predictions are quite consistent, since our model has nearly the same count for 0 and 1 than the target variable in our data set.

```
confusion_matrix <- table(diabetes_data$YBin, prediction_1se)
confusion_matrix
```

```
      prediction_1se
      0    1
0 166  55
1  54 167
```

```
accuracy <- (confusion_matrix[1,1] + confusion_matrix[2,2]) / nrow(diabetes_data)
accuracy
```

```
[1] 0.7533937
```

The accuracy is equal to 76.6 %.

K-Fold function

Let's code a function that will take not only the number of folds but also the lambda so we can use it for both values of λ_{min} and λ_{1se} in our ridge regression model.

```
kfold_ridge <- function(k, lambda)
{
  # create a vector of length number of folds
  performance <- vector(length = k)

  # create a sequence from 1 to k
  folds <- cut(seq(1,nrow(diabetes_data)), breaks = k, labels = FALSE)

  # perform 10 fold cross validation
  for(i in 1:k)
  {
    # split data by fold
    index <- which(folds == i, arr.ind = TRUE)
    test_data <- diabetes_data[index,]
    train_data <- diabetes_data[-index,]

    # train the logistic regression on the train data set
    reg_log <- glmnet(x = train_data[,-11], y = train_data$YBin,
                      alpha = 0, family = "binomial", lambda = lambda)


    # make predictions
    prediction <- as.numeric(
      predict(reg_log, as.matrix(test_data[,-11]), type = "response") > 0.5)
    confusion_matrix <- table(prediction, test_data$YBin)

    # compute the performance
    performance[i] <- (confusion_matrix[1,1] + confusion_matrix[2,2]) / nrow(test_data)
  }

  # returning the vector of performances
  return (performance)
}
```


K-Fold boxplots

```
boxplot(kfold_ridge(10, lambda_min), kfold_ridge(10, lambda_1se), main = "K-Fold",
        names = c("lambda_min", "lambda_1se"), col = c("royalblue1", "tomato1"))
```




In our case we did two 10-Folds with λ_{min} and λ_{1se} . By looking at the boxplots we can see that the median of our accuracy is between 75% and 80% for the two models. But we can clearly observe that the box for λ_{1se} is higher than the one for λ_{min} which means in general we will have a better accuracy by using λ_{1se} . More than that, we know that with λ_{1se} we are having a larger penalization than with λ_{min} which explains why we have a larger interquartile range and a larger distance between the minimum and the maximum.

Lasso Regression

Regularization Path

```
lasso <- glmnet(x = train_data[, -11], y = train_data$YBin, alpha = 1, family = "binomial")  
plot(lasso, xvar = "lambda", label = TRUE, lwd = 2)
```




In the plot above, we can see the evolution of the coefficient values depending on λ . Knowing that we are doing a lasso regression, we can tell that they will all converge to 0 one by one and not all at the same time.

When a data set has many co-variables, lasso can be used to identify and extract those co-variables which are the most significant. Switching to the lasso penalty also conducts automated variable selection.

Cross-Validation

```
cv_lasso <- cv.glmnet(as.matrix(train_data[,-11]), train_data$YBin,  
                      family = "binomial", alpha = 1, type.measure = "mse")  
plot(cv_lasso)
```



We observe a slight improvement in the Mean Squared Error as our penalty $\log(\lambda)$ gets larger, suggesting that a regular OLS model likely overfits the training data. But as we constrain it further by continuing to increase the penalty our Mean Squared Error starts to increase.

Lambda Min

```
lambda_min <- cv_lasso$lambda.min
lasso_min <- glmnet(x = train_data[,-11], y = train_data$YBin,
                   alpha = 1, family = "binomial", lambda = lambda_min)
lasso_min$beta
```

```
10 x 1 sparse Matrix of class "dgCMatrix"
      s0
AGE  0.011199543
SEX -1.186277595
BMI  0.158657128
BP   0.036306767
S1  -0.007128551
S2   .
S3  -0.036311964
S4   .
S5   1.652994086
S6   0.001486600
```

The model obtained using λ_{min} gives us the model with lowest Mean Squared Error, which can seem to be a good choice but in fact it can implies overfitting. We observe that the model does not contain *AGE*, *S2* and *S4*.

```
prediction_min <- as.numeric(
  predict(lasso_min, as.matrix(diabetes_data[,-11]), type = "response") > 0.5)
```

Using the Maximum A Posteriori criteria we can make predictions for our binary variable *YBin*:

```
table(prediction_min)
```

```
prediction_min
 0    1
219 223
```

By comparing with the table of the target variable we can tell that our predictions are quite consistent, since our model has nearly the same count for 0 and 1 than the target variable in our data set.

```
confusion_matrix <- table(diabetes_data$YBin, prediction_min)
confusion_matrix
```

```
      prediction_min
      0    1
0 163  58
1  56 165
```

```
accuracy <- (confusion_matrix[1,1] + confusion_matrix[2,2]) / nrow(diabetes_data)
accuracy
```

```
[1] 0.7420814
```

The accuracy is equal to 76.5 %.

Lambda 1se

```
lambda_1se <- cv_lasso$lambda.1se
lasso_1se <- glmnet(x = train_data[,-11], y = train_data$YBin,
                    alpha = 1, family = "binomial", lambda = lambda_1se)
lasso_1se$beta
```

```
10 x 1 sparse Matrix of class "dgCMatrix"
      s0
AGE    .
SEX -0.015962142
BMI  0.116206311
BP   0.015752329
S1    .
S2    .
S3 -0.008381972
S4    .
S5  0.997527620
S6    .
```

The model obtained using λ_{1se} gives us a simpler model which can avoid overfitting. We observe that the model does not contain *AGE*, *S1*, *S2*, *S4* and *S6*, so the model is simpler than the one with λ_{min} .

```
prediction_1se <- as.numeric(
  predict(lasso_1se, as.matrix(diabetes_data[,-11]), type = "response") > 0.5)
```

Using the Maximum A Posteriori criteria we can make predictions for our binary variable *YBin*:

```
table(prediction_1se)
```

```
prediction_1se
 0    1
222 220
```

By comparing with the table of the target variable we can tell that our predictions are quite consistent, since our model has nearly the same count for 0 and 1 than the target variable in our data set.

```
confusion_matrix <- table(diabetes_data$YBin, prediction_1se)
confusion_matrix
```

```
      prediction_1se
      0    1
0 165  56
1  57 164
```

```
accuracy <- (confusion_matrix[1,1] + confusion_matrix[2,2]) / nrow(diabetes_data)
accuracy
```

```
[1] 0.7443439
```

The accuracy is equal to 74.8%.

K-Fold function

Let's code a function that will take not only the number of folds but also the lambda so we can use it for both values of λ_{min} and λ_{1se} in our ridge lasso model.

```
kfold_lasso <- function(k, lambda)
{
  # create a vector of length number of folds
  performance <- vector(length = k)

  # create a sequence from 1 to k
  folds <- cut(seq(1,nrow(diabetes_data)), breaks = k, labels = FALSE)

  # perform 10 fold cross validation
  for(i in 1:k)
  {
    # split data by fold
    index <- which(folds == i, arr.ind = TRUE)
    test_data <- diabetes_data[index,]
    train_data <- diabetes_data[-index,]

    # train the logistic regression on the train data set
    reg_log <- glmnet(x = train_data[,-11], y = train_data$YBin,
                      alpha = 1, family = "binomial", lambda = lambda)

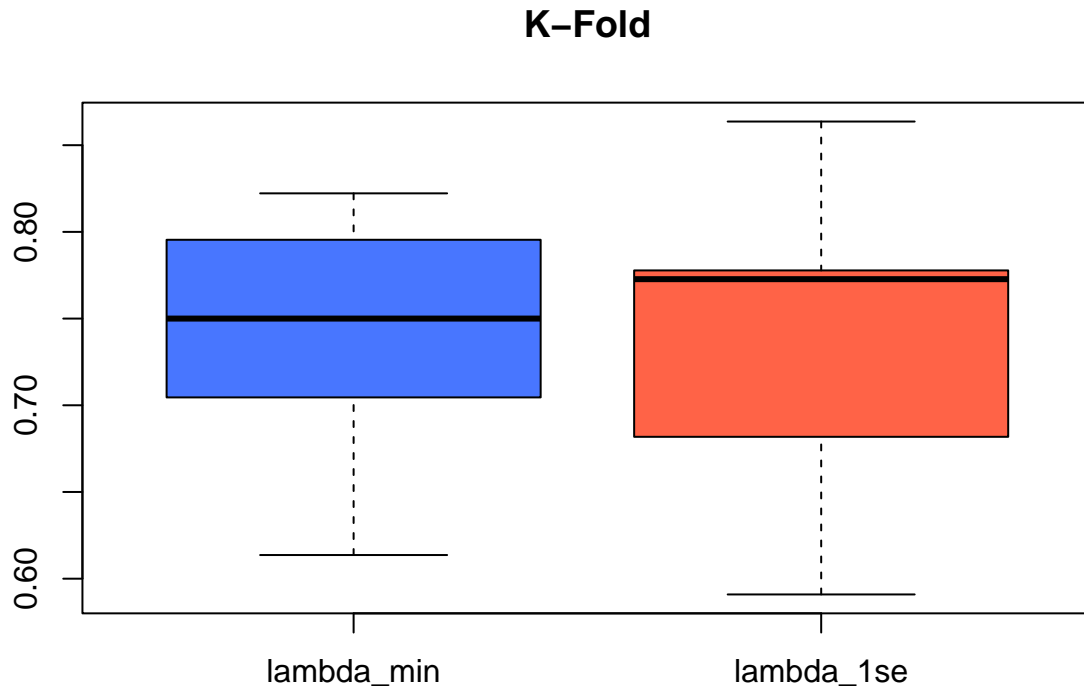
    # make predictions
    prediction <- as.numeric(
      predict(reg_log, as.matrix(test_data[,-11]), type = "response") > 0.5)

    confusion_matrix <- table(prediction, test_data$YBin)
    # compute the performance
    performance[i] <- (confusion_matrix[1,1] + confusion_matrix[2,2]) / nrow(test_data)
  }
  # returning the vector of performances
  return (performance)
}
```

We could code a function that takes not only the number of folds and the lambda but also the alpha which will be equal to 0 for a ridge regression and 1 for a lasso regression.

K-Fold boxplots

```
boxplot(kfold_lasso(10, lambda_min), kfold_lasso(10, lambda_1se), main = "K-Fold",
        names = c("lambda_min", "lambda_1se"), col = c("royalblue1", "tomato1"))
```



In our case we did two 10-Folds with λ_{min} and λ_{1se} . By looking at the boxplots we can see that the median of our accuracy is between 75% and 80% for the two models. But we can clearly observe that the box for λ_{1se} is higher than the one for λ_{min} which means in general we will have a better accuracy by using λ_{1se} . More than that, we know that with λ_{1se} we are having a larger penalization than with λ_{min} which explains why we have a larger interquartile range and a larger distance between the minimum and the maximum.

Conclusion

Let's compare the accuracy of all the models obtained using a K-Fold procedure for Cross Validation.

In order to do that let's code a function to make k-fold on the null and the step-wise model.

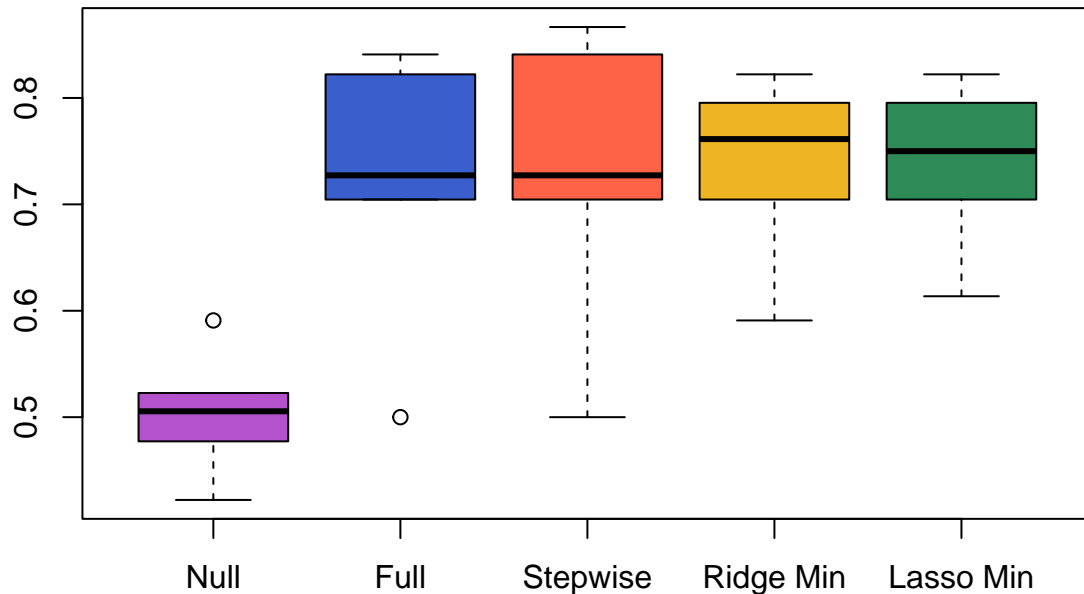
```
kfold_none <- function(k) # k is here the number of folds
{
  # create a vector of length number of folds
  performance <- vector(length = k)
  # create a sequence from 1 to k
  folds <- cut(seq(1,nrow(diabetes_data)), breaks = k, labels = FALSE)
  # perform k fold cross validation
  for(i in 1:k)
  {
    # split data by fold
    index <- which(folds == i, arr.ind = TRUE)
    test_data <- diabetes_data[index,]
    train_data <- diabetes_data[-index,]
    # train the logistic regression on the train data set
    reg_log <- glm(YBin ~ 1, family = binomial, data = train_data)
    # make predictions
    prediction <- predict(reg_log, test_data)
    # compute confusion matrix and performance
    confusion_matrix <- table(as.numeric(prediction > 0.5), diabetes_data[index,]$YBin)
    performance[i] <- confusion_matrix[1,2] / nrow(test_data)
  }
  return (performance)
}
```

```
kfold_step <- function(k) # k is here the number of folds
{
  # create a vector of length number of folds
  performance <- vector(length = k)
  # create a sequence from 1 to k
  folds <- cut(seq(1,nrow(diabetes_data)), breaks = k, labels = FALSE)
  # perform k fold cross validation
  for(i in 1:k)
  {
    # split data by fold
    index <- which(folds == i, arr.ind = TRUE)
    test_data <- diabetes_data[index,]
    train_data <- diabetes_data[-index,]
    # train the logistic regression on the train data set
    reg_log <- step(glm(YBin ~ ., family = binomial, data = train_data),
                    direction = "both", trace = FALSE)
    # make predictions
    prediction <- predict(reg_log, test_data)
    # compute confusion matrix and performance
    confusion_matrix <- table(as.numeric(prediction > 0.5), diabetes_data[index,]$YBin)
    performance[i] <- (confusion_matrix[1,1] + confusion_matrix[2,2]) / nrow(test_data)
  }
  return (performance)
}
```



```
boxplot(kfold_none(10), kfold_all(10), kfold_step(10),
       kfold_ridge(10, cv_ridge$lambda.min), kfold_lasso(10, cv_lasso$lambda.min),
       main = "10-Fold for 5 different models",
       names = c("Null", "Full", "Stepwise", "Ridge Min", "Lasso Min"),
       col = c("mediumorchid3", "royalblue3", "tomato1", "goldenrod2", "seagreen"))
```

10-Fold for 5 different models



By looking at the boxplots we can compare between all the models we did in this practical work:

- We can notice the presence of outliers;
- The Null model has the lowest accuracy because he doesn't take in count any co-variables;
- The Stepwise model has a very small minimum value which is near to the median of the Null model;
- The Full and Stepwise models have a better accuracy than the Null model but their median are lower than the median for Ridge and Lasso, we can explain that by the fact that Ridge and Lasso are applying penalization in order to avoid overfitting so it has a better predictive power on new data set which is the case when doing K-Fold Cross Validation;
- For Ridge and Lasso we can clearly see that they have the best accuracy but we can observe than the interquartile range and the distance between min and max values are larger for Ridge than for Lasso because Lasso does variable selection. Generally, when we have many small or medium sized effects we should go with Ridge. If we have only a few variables with a medium or large effect, we should go with Lasso;
- "Ridge regression does a proportional shrinkage. Lasso translates each coefficient by a constant factor, truncating at zero." from **The Elements of Statistical Learning: Data Mining, Inference, and Prediction. Hastie, Tibshirani, Friedman.**